



Learning Machine Learning with Kaggle Challenges

(3) Deep Learning



Qiyang Hu
IDRE

Interest over time from Google Trends




[2018 ACM Turing Award Recipients](#)

A logistic-regression classifier ~ one artificial neuron




(Deep) Neural Networks ~ piling/stacking logistic-regression classifiers



How deep a deep learning network can be?



- LeNet-5 (1998)



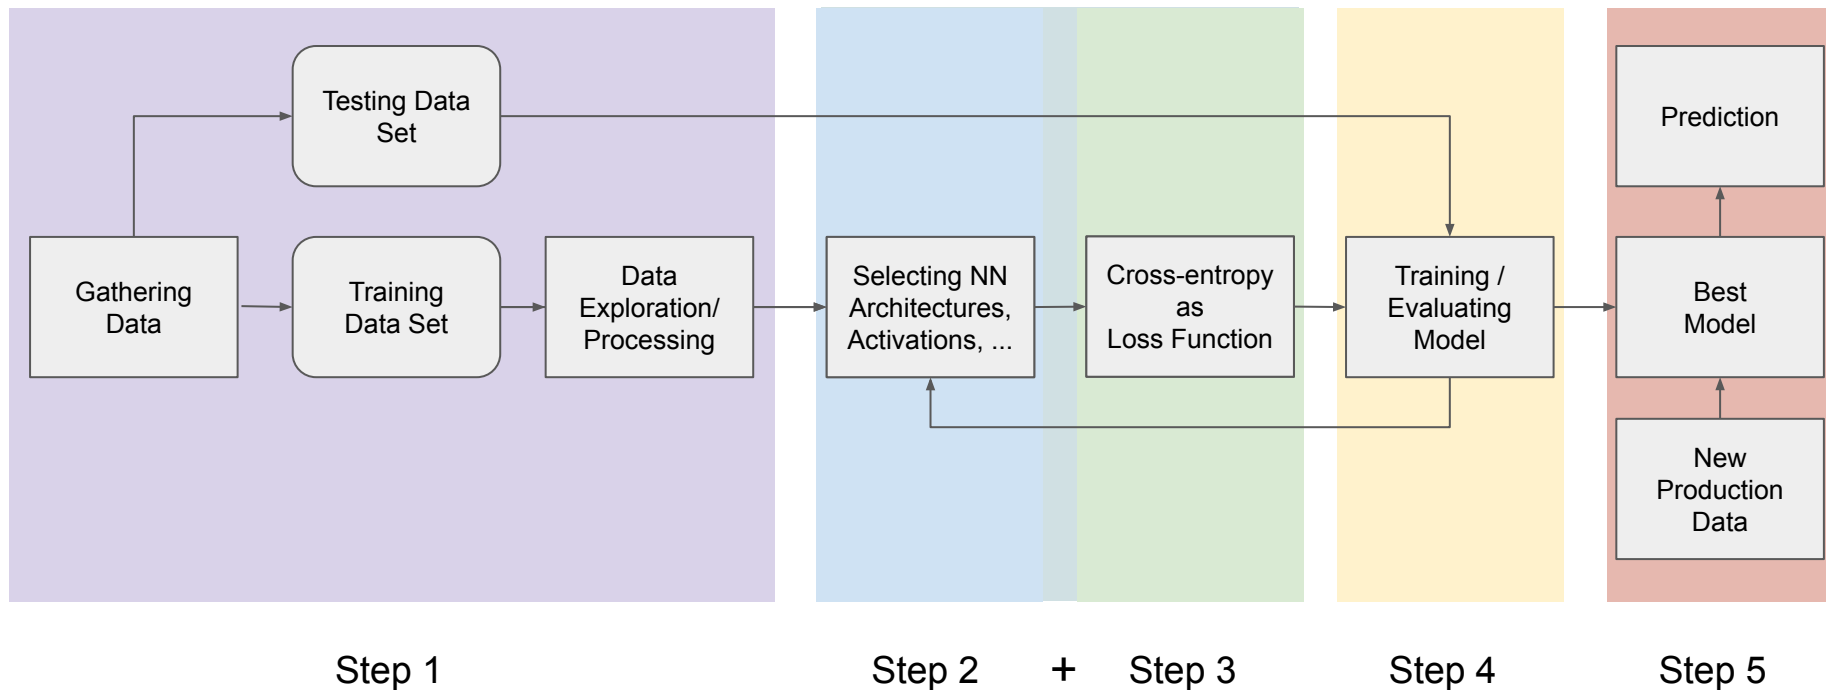
Year	CNN	Developed by	Place	Top-5 error rate	No. of parameters
1998	LeNet(8)	Yann LeCun et al			60 thousand
2012	AlexNet(7)	Alex Krizhevsky, Geoffrey Hinton, Ilya Sutskever	1st	15.3%	60 million
2013	ZFNet()	Matthew Zeiler and Rob Fergus	1st	14.8%	
2014	GoogLeNet(19)	Google	1st	6.67%	4 million
2014	VGG Net(16)	Simonyan, Zisserman	2nd	7.3%	138 million
2015	ResNet(152)	Kaiming He	1st	3.6%	

Why deep?

- Shallow network can fit any function
 - Has less number of hidden layers
 - Has to be really “fat”
- Deep network is more efficient.
 - It can extract/build better features
 - Exponentially fewer parameters ([2017](#))




Workflow for a **deep** learning project



Types of Neural Network Architectures


- Feed forward neural networks (No cycle in node connections)

- Perceptron
- Fully connected network
- Convolutional networks



- Recurrent networks (w/ directed cycle in node connections)


- Fully recurrent NN
- Recursive NN
- Long short-term memory (LSTM)
- Symmetrically connected networks
 - Hopfield network (w/o hidden nodes)
 - Deep Boltzmann Machine (w/ hidden nodes)



Activation Function

- Sigmoid function:

$$\sigma(z) = \frac{1}{1 + \exp(-z)}$$



- Rectified linear unit (ReLU)

$$f(x) = x^+ = \max(0, x)$$


- Softplus
- Leaky ReLU
- Exponential LU (ELUs)

- Softmax function:

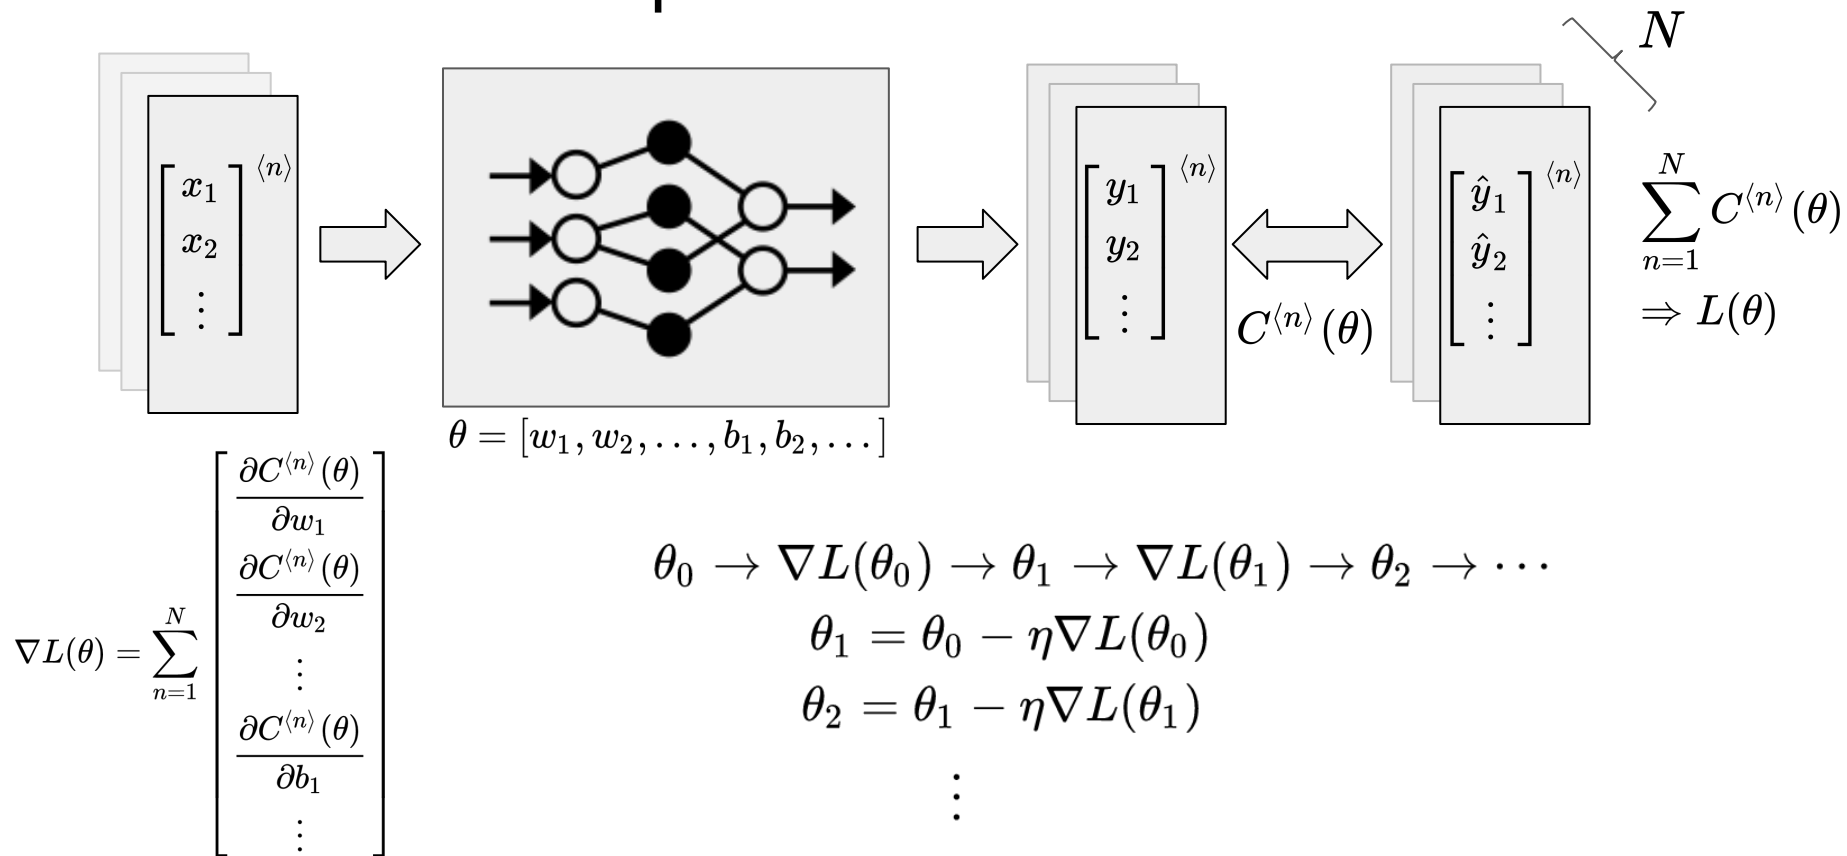
$$y_i = \frac{e^{z^{(i)}}}{\sum_{j=0}^K e^{z^{(j)}}}$$

- Maxout Network:


- *Learnable* activation function



How to train a deep neural network?



Backpropagation: a game of chain rule



$$y = \sigma_L \left(w_L \cdot \sigma_{L-1} \left(\cdots w_2 \cdot \sigma_1 \left(\underbrace{w_1 \cdot x + b_1}_{z_1} \right) + b_2 \right) + b_L \right)$$

$$\frac{\partial C(y(w) - \hat{y})}{\partial w} = \frac{\partial z}{\partial w} \frac{\partial C}{\partial z} = \frac{\partial z}{\partial w} \left[\frac{\partial a}{\partial z} \frac{\partial C}{\partial a} \right] = \frac{\partial z}{\partial w} \left[\sigma' \cdot \left(\underbrace{\frac{\partial z_{(+1)}}{\partial a} \frac{\partial C}{\partial z_{(+1)}}}_{a_1} \right) \right]$$

① Forward Pass

$$\frac{\partial z_1}{\partial w_1} = x_1 \longrightarrow \frac{\partial z_2}{\partial w_2} = a_1 \longrightarrow \cdots \longrightarrow \frac{\partial z_{L-1}}{\partial w_{L-1}} = a_{L-2} \longrightarrow \frac{\partial z_L}{\partial w_L} = a_{L-1}$$

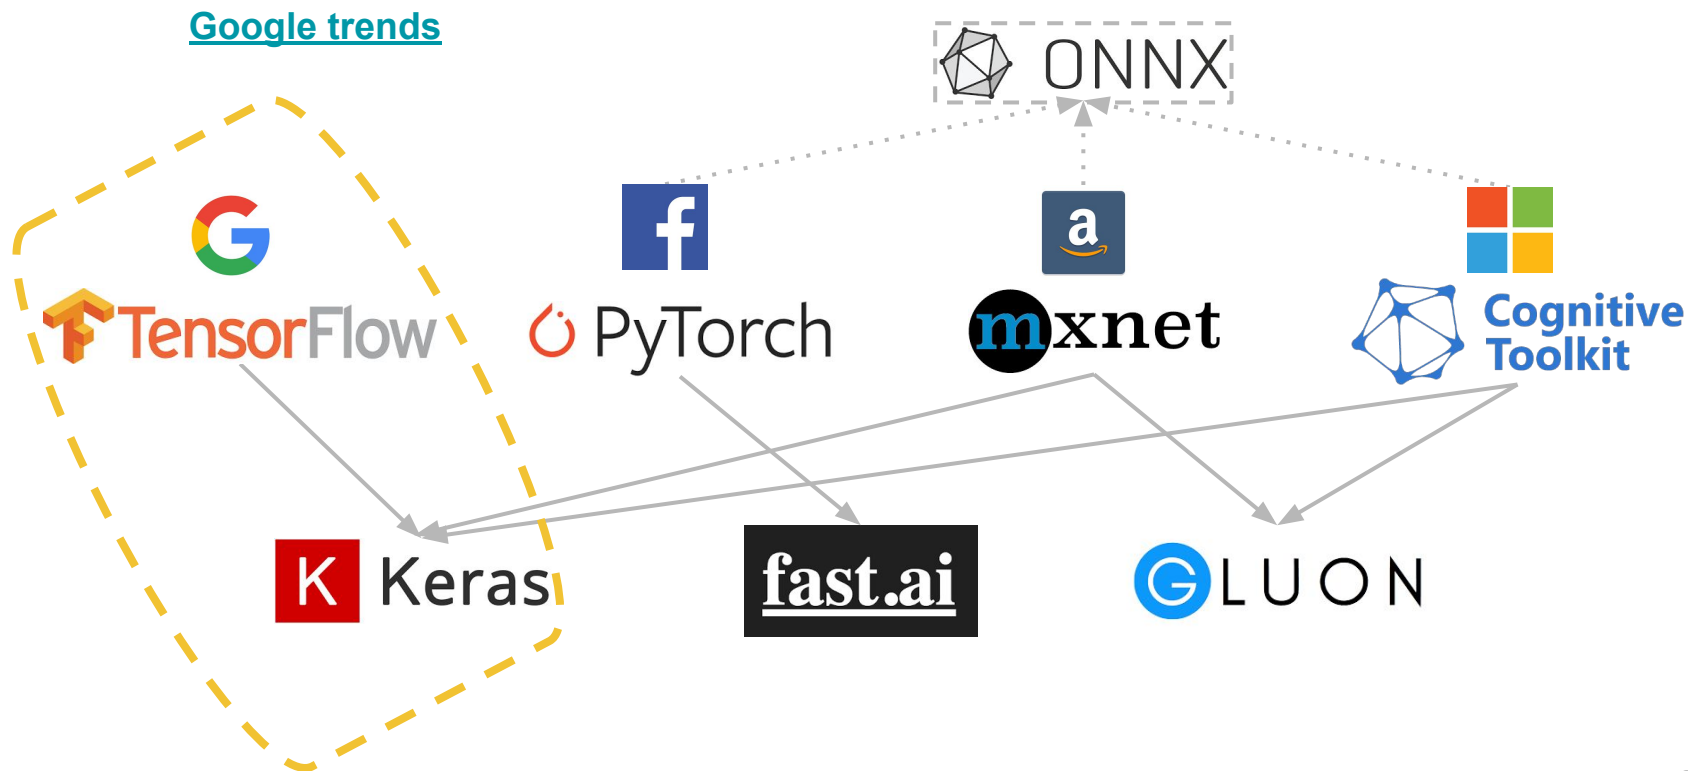
② Backward Pass

$$\frac{\partial C}{\partial z_1} = \sigma'_1 \left[w_2 \frac{\partial C}{\partial z_2} \right] \longleftarrow \cdots \longleftarrow \frac{\partial C}{\partial z_{L-1}} = \sigma'_{L-1} \left[w_L \frac{\partial C}{\partial z_L} \right] \longleftarrow \frac{\partial C}{\partial z_L} = \sigma'_L \frac{\partial C}{\partial y} \longleftarrow \frac{\partial C}{\partial y}$$

Deep learning frameworks



Mainstream Players



Tensorflow v2.0 just out on Sep. 30, 2019

- **BIG** Changes from v1.x

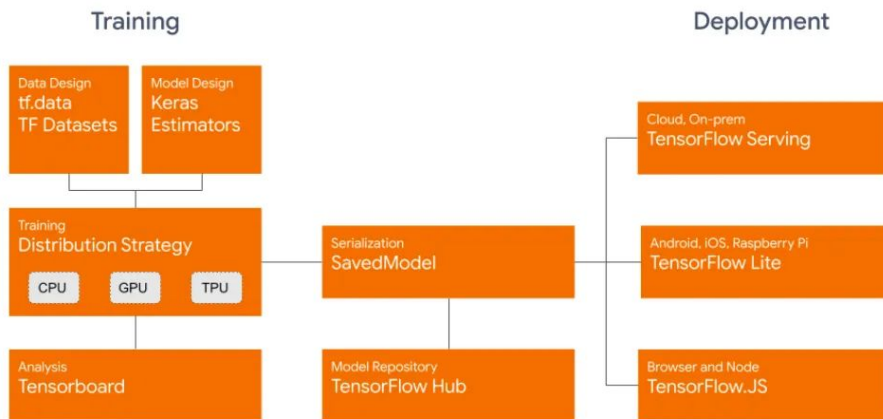
- Keras as a core API
- Eager execution by default
- `tf.function` decorator to speed up
- `tf.data` to build complex input pipelines
- Model deployment to various platforms

- Cautions

- Poor compatibility with v1.x
- Lots of confusions from API names

- Importing tf2.0 in standalone python

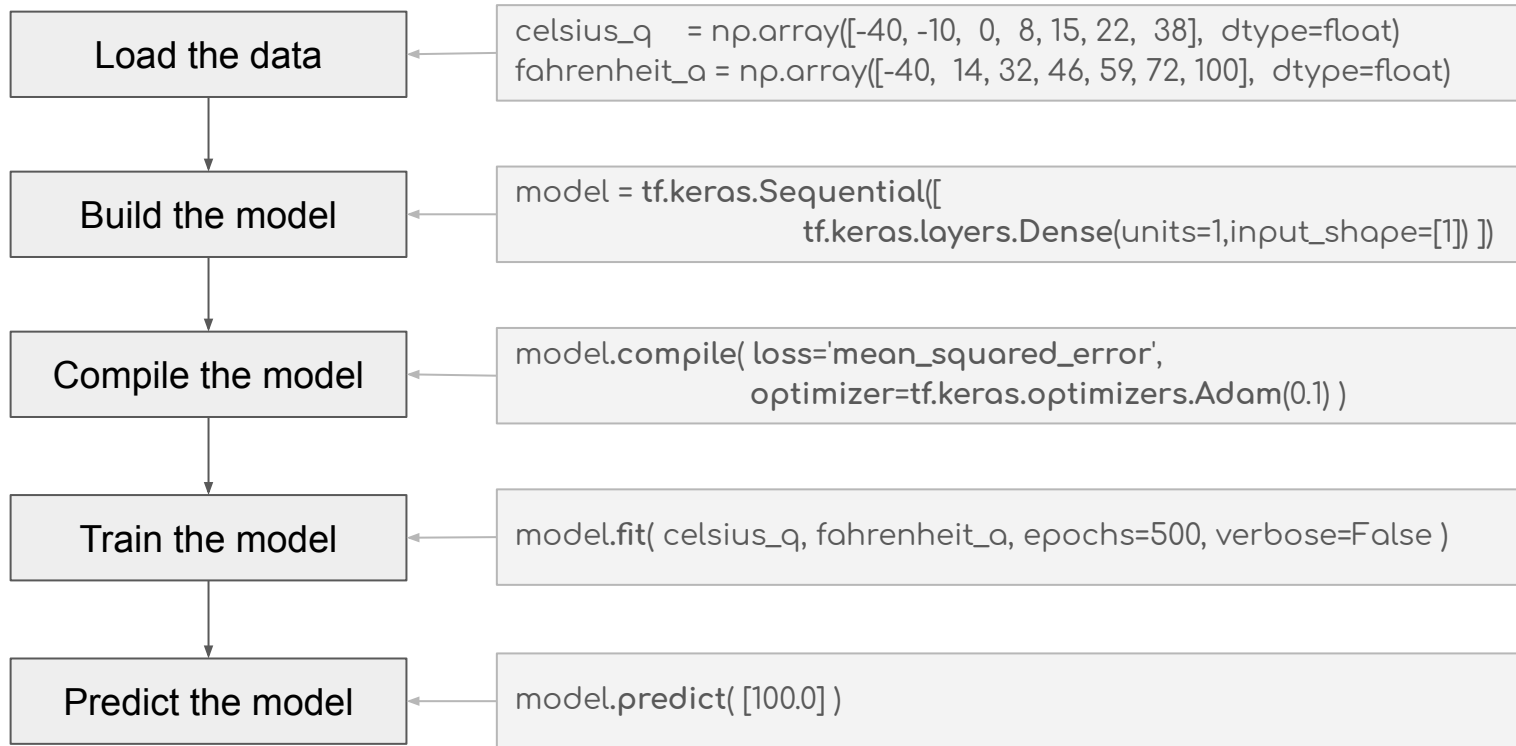
- `import tensorflow.compat.v2 as tf`
`tf.enable_v2_behavior()`



- Importing tf2.0 in Colab

- `%tensorflow_version 2.x`

Using tf.keras [\(an 1-neuron 1-layer example\)](#)




Convolutional Neural Networks (CNNs)


- A special network architecture to reduce parameters

- 3 processes in CNNs:

- Convolutions to extract as tiles



- Poolings to downsample



- Flattening

- Logic behind CNNs

- Sparse connectivity (characteristic features in smaller local regions)
- Parameter equivariance & sharing (features appear in different locations)
- Translation invariance (some sampling will not lose main information)

One Channel, One Filter

0	0	0	0	0	0
0	105	102	100	97	96
0	103	99	103	101	102
0	101	98	104	102	100
0	99	101	106	104	99
0	104	104	104	100	98

Image Matrix

Kernel Matrix		
0	-1	0
-1	5	-1
0	-1	0

320				

Output Matrix

$$\begin{aligned} &0 * 0 + 0 * -1 + 0 * 0 \\ &+ 0 * -1 + 105 * 5 + 102 * -1 \\ &+ 0 * 0 + 103 * -1 + 99 * 0 = 320 \end{aligned}$$

**Convolution with horizontal and
vertical strides = 1, with 'same' padding**

Multiple Channels

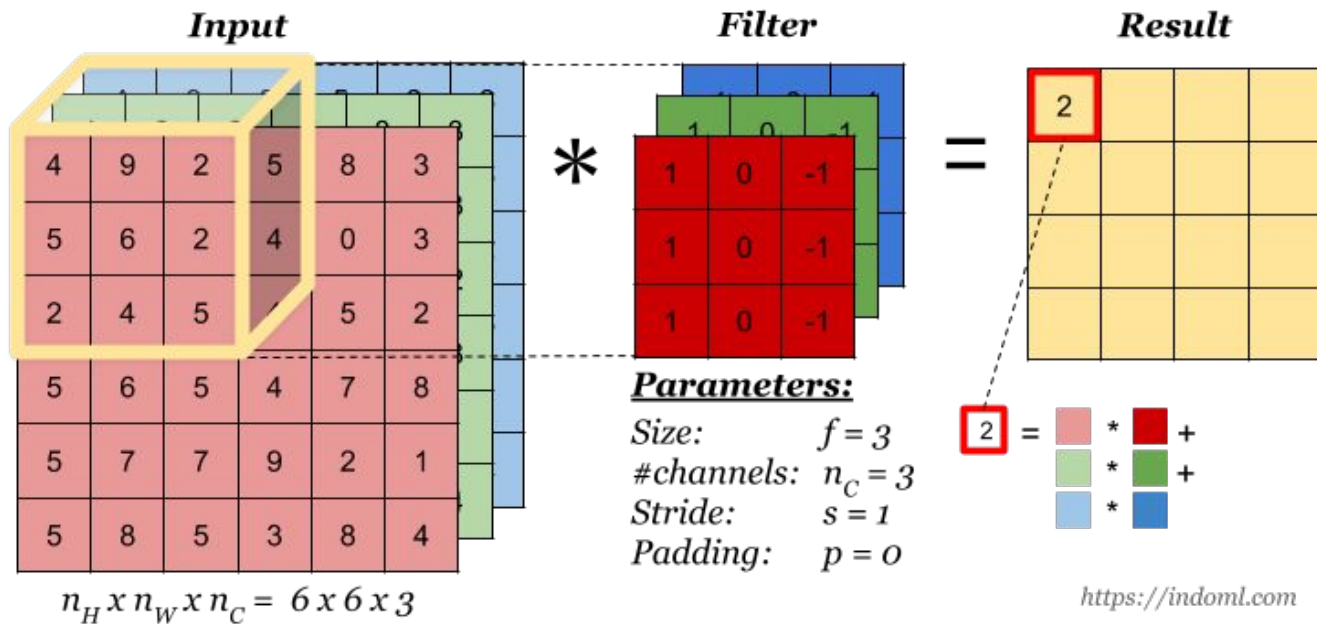
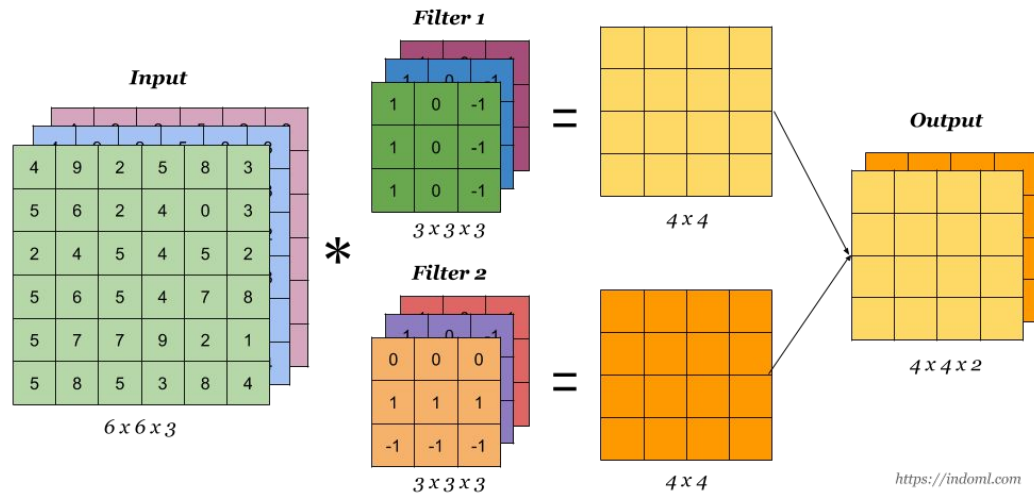
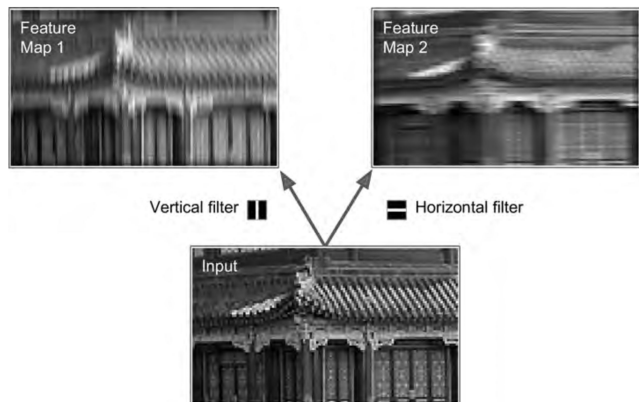


Figure [Source](#)

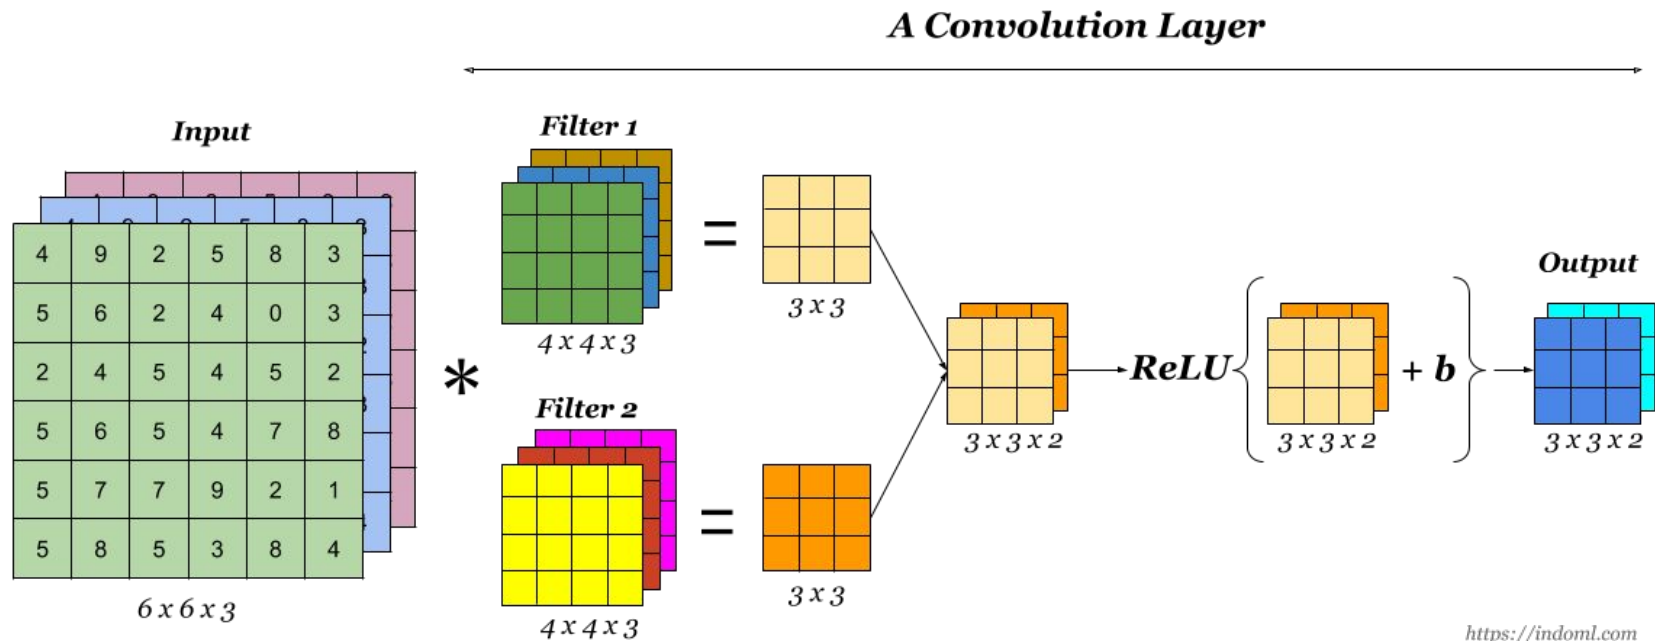
Multiple Filters



Figures from Aurélien Géron's [1st Ed. Book](#)

Figure [Source](#)

A Convolutional layer



<https://indoml.com>

Figure [Source](#)

Pooling Layer

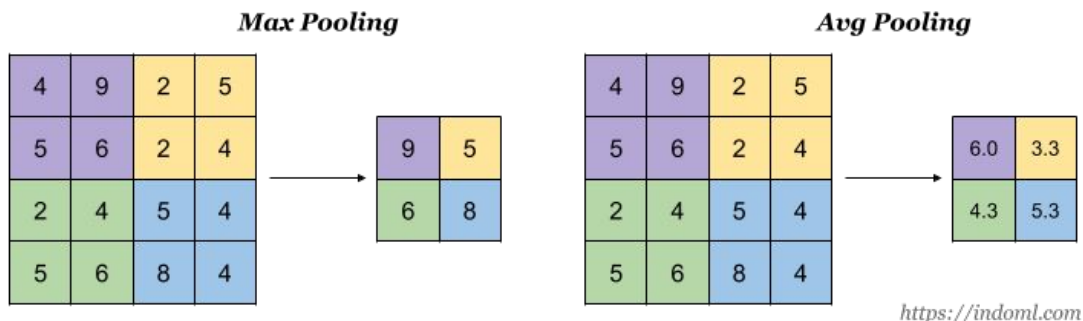



Figure [Source](#)

- Assuming downsampling will not lose the major information.



Figures from Aurélien Géron's [1st Ed. Book](#)

Architecture of Convolutional Neural Networks

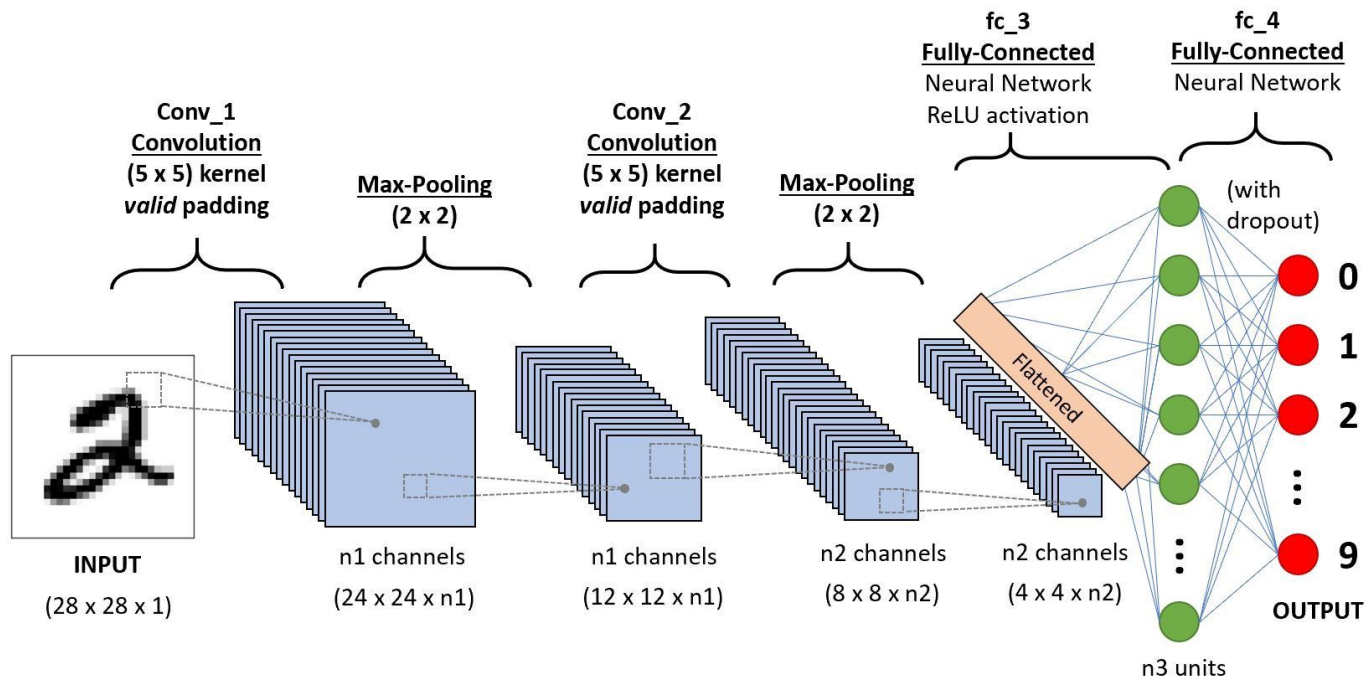


Figure [Source](#)

Dogs vs. Cats Kaggle Challenge

- Redux: Kernels Edition

- Submission scored by the probability of dogs using log loss

$$L = -\frac{1}{n} \sum_{i=1}^n \left[y_i \log(\hat{y}_i) + (1 - y_i) \log(1 - \hat{y}_i) \right]$$

- Dataset

- Training set: 25,000 dogs and cats images
- Testing set: 12,500 images


- Two concerns:

- Images with different sizes
 - Neural network needs fixed sized input.
 - We will resize images to 150x150 pixels
- Images are colored
 - Represented by Red-Green-Blue channels
 - One image \Rightarrow 150x150x3 matrices



Import image data using tf.keras

- tf.keras.preprocessing.image.ImageDataGenerator.flow_from_directory



```
train_img_gen = ImageDataGenerator(rescale=1./255)
train_data_gen = train_img_gen.flow_from_directory(batch_size=BATCH_SIZE,
                                                    directory=train_dir,
                                                    shuffle=True,
                                                    target_size=(IMG_SHAPE,IMG_SHAPE), #(150,150)
                                                    class_mode='binary')
```


Construct CNN architecture using tf.keras

- Convolution layer:

[tf.keras.layers.Conv2D\(filters, kernel_size, activation, ...\)](#)

- Filters (feature maps): 32 or 64 or 128 ...
- Kernel_size: (3,3)
- Activation function: 'relu'
- Input_shape: (150, 150, 3)

- MaxPooling layer:

[tf.keras.layers.MaxPooling2D\(pool_size, strides, ...\)](#)

- Pool_size: 2
- Strides: 2

- Flattened layer

[tf.keras.layers.Flatten\(\)](#)

- Dense layer

[tf.keras.layers.Dense\(units, activation, ...\)](#)

- Units: 512 and 2
- Activation: 'relu' and 'softmax'

4 Conv blocks with maxpool in each

```
model = tf.keras.models.Sequential([  
    tf.keras.layers.Conv2D(32, (3,3), activation='relu',  
                           input_shape=(150, 150, 3)),  
    tf.keras.layers.MaxPooling2D(2, 2),  
    tf.keras.layers.Conv2D(64, (3,3), activation='relu'),  
    tf.keras.layers.MaxPooling2D(2,2),  
    tf.keras.layers.Conv2D(128, (3,3), activation='relu'),  
    tf.keras.layers.MaxPooling2D(2,2),  
    tf.keras.layers.Conv2D(128, (3,3), activation='relu'),  
    tf.keras.layers.MaxPooling2D(2,2),  
    tf.keras.layers.Flatten(),  
    tf.keras.layers.Dense(512, activation='relu'),  
    tf.keras.layers.Dense(2, activation='softmax')  
])
```

Don't forget to

- Sign in your info to the class
 - To get the email notifications
- Contact me for questions or discussions
 - hugy@idre.ucla.edu
 - Office: Math Sci #3330
 - Phone: 310-825-2011
- Fill out the survey for comments:
 - <https://forms.gle/t3f8CztFQpeFFksy6>

